

# Design Review of a Nova + CycleFold Implementation: Security and Scaling Considerations for Recursive SNARK Systems

Vishal Singh  
ZipperLabs

January 28, 2026

## Abstract

Recursive proof systems allow a prover to produce a succinct proof of an arbitrarily long computation by repeatedly folding (or aggregating) proofs and re-proving the verification of the previous step. Nova [1] is a prominent design in this space: it enables efficient incremental verifiable computation (IVC) via folding schemes over relaxed R1CS. CycleFold [2] addresses practical challenges around recursion on curve cycles.

This paper is written for security researchers and protocol engineers. We present a design-review-oriented description of a concrete “Nova + CycleFold” implementation, with the goal of making its security and performance assumptions explicit. We provide an end-to-end component breakdown (circuits, commitments, transcripts, and accumulator updates), enumerate relevant threat models (soundness, zero-knowledge, malleability, and side channels), and give a checklist of invariants and test vectors to support independent review.

## 1 Introduction

Succinct non-interactive arguments of knowledge (SNARKs) have become a foundational primitive for privacy-preserving computation and blockchain scalability. While single-shot SNARKs can attest to the correct execution of a fixed computation, many real systems require *incremental* proofs over long-running processes: state transitions of a ledger, continuous monitoring, or streaming data processing. Recursive proof systems address this need by proving that “the next step was computed correctly *and* the previous proof verified.”

Nova introduced an efficient approach to incremental verifiable computation (IVC) using a *folding scheme* over *relaxed* Rank-1 Constraint Systems (R1CS) [1]. In this paper we focus on an implementation that combines Nova’s folding scheme with CycleFold [2], a complementary technique intended to improve recursion ergonomics when working with curve-cycle constraints.

This paper is a *design review aid*. We assume the reader is familiar with finite fields, elliptic curves, and basic SNARK concepts, but may not have prior exposure to Nova-style folding.

Our contributions are:

- a component-level breakdown of a Nova + CycleFold implementation, including precise interface boundaries;
- a security-focused discussion of transcripts, commitments, public input binding, and malleability risks;
- a review checklist (invariants, negative tests, and interoperability/test-vector requirements);
- scaling considerations and implementation pitfalls that commonly affect correctness or auditability.

## 2 Background and Problem Setting

### 2.1 R1CS and Arithmetization

An R1CS instance over a field  $F$  consists of matrices  $(A, B, C)$  such that for a witness vector  $z$ ,

$$(Az) \circ (Bz) = Cz, \quad (1)$$

where  $\circ$  is the Hadamard (entrywise) product. Implementation choices include: (i) representing constraints densely or sparsely, (ii) selecting an arithmetization (R1CS vs. PLONKish arithmetics), and (iii) deciding how to express hashing, Merkle paths, and elliptic-curve operations.

### 2.2 Incremental Verifiable Computation (IVC)

Let  $\text{Step}$  be a transition function that maps a state  $s_i$  and input  $x_i$  to a new state  $s_{i+1}$ . IVC aims to produce a succinct proof  $\pi_T$  that  $s_T$  is the result of applying  $\text{Step}$  for  $T$  steps, without producing a proof whose size grows linearly in  $T$ .

## 3 Nova in a Nutshell

Nova’s core idea is a *folding scheme* for relaxed R1CS instances. Intuitively, instead of proving each step independently, the prover maintains an accumulator instance that represents “all prior work” and updates it each step. A final SNARK then proves that the accumulator is valid.

### 3.1 Relaxed R1CS

Relaxed R1CS introduces a slack term to enable linear combination of instances. A common representation is:

$$(Az) \circ (Bz) = Cz + u \cdot e, \quad (2)$$

where  $u \in F$  and  $e$  is an “error” vector. Folding combines two instances into one by sampling a challenge  $r$  and forming linear combinations of witnesses and residuals.

### 3.2 Recursion-Friendly Curves and Cycle Constraints

Implementations typically require a curve cycle (or compatible curves) so that the verifier circuit for one proof system can be expressed over the scalar field of another. CycleFold [2] targets this pain point by providing a folding-based pathway for composing proofs across curve cycles with a focus on implementability. Engineering constraints include:

- availability of optimized field arithmetic for both fields;
- endomorphisms and windowing strategies for MSMs (multi-scalar multiplications);
- serialization formats and cross-language bindings.

## 4 Privacy: Threat Models and Design Choices

### 4.1 What “Privacy” Means Here

We distinguish between:

1. **Witness privacy (zero-knowledge):** the proof reveals nothing about private inputs beyond validity.
2. **State privacy:** the committed state hides user balances/records while allowing valid transitions.
3. **Metadata privacy:** hiding access patterns, step counts, or timing information.

### 4.2 Zero-Knowledge in Nova-based Systems

Nova itself is an IVC framework; privacy depends on the SNARK used for the final proof and the commitments used inside folding. Implementation pitfalls:

- **Transcript hygiene:** domain separation and consistent Fiat–Shamir transcripts across languages.
- **Randomness generation:** secure RNGs and deterministic modes for reproducibility.
- **Side channels:** variable-time field operations, cache leakage in MSMs, and branching on secret data.

### 4.3 Commitment Schemes and Hashing

Nova-style designs rely on commitments to witnesses/instances. In practice, choices are dominated by circuit cost:

- **Hash in-circuit:** Poseidon/Rescue-type hashes (algebraic) vs. Keccak/SHA (bit-based, typically expensive in arithmetic circuits).
- **Commitment type:** Pedersen-style commitments (EC-based) vs. polynomial commitments depending on the outer SNARK.

## 5 Scaling: Engineering for Throughput and Latency

### 5.1 The Scaling Objective

Scaling goals vary:

- **Low-latency proofs:** fast per-block proving for rollups.
- **High-throughput batching:** aggregate many user actions per step.
- **Long horizon computation:** prove months of history with bounded resources.

### 5.2 Bottlenecks and Cost Centers

In typical implementations, the dominant costs are:

1. MSMs and FFT-like operations (if the outer SNARK uses polynomial commitments),
2. hashing/Merkle path verification for state commitments,
3. memory bandwidth from large witnesses and constraint matrices,
4. cross-thread synchronization in parallel proving.

### 5.3 Batching Patterns

We outline three common patterns:

**Transaction batching:** represent  $k$  transactions per step; prove signature checks and state updates in one circuit.

**Proof batching:** fold multiple independent accumulator updates into one step via tree-style aggregation.

**Data availability batching:** commit to large blobs (e.g., calldata) and prove consistency with state transitions.

## 6 Implementation Blueprint

### 6.1 Reference Architecture

A practical Nova deployment typically separates:

1. **Step circuit:** enforces  $s_{i+1} = \text{Step}(s_i, x_i)$  and exposes a commitment to  $s_{i+1}$ .
2. **Folding layer:** updates the relaxed R1CS accumulator with a Fiat–Shamir challenge.
3. **Final SNARK:** proves the folded accumulator is valid and binds to the final public output.

### 6.2 Interface Contracts

We recommend explicit, versioned interfaces:

- canonical encoding for field elements and curve points;
- fixed transcript labels for every challenge;
- stable public input ordering (breaks here invalidate verifiers).

### 6.3 Practical Optimizations

Common optimizations include:

- **Constraint minimization:** custom gates for hashing and EC operations when available.
- **Witness streaming:** avoid materializing entire witnesses when the system supports chunked constraints.
- **Parallel MSM:** split scalars across threads with pinned memory and NUMA-aware allocation.
- **Amortized setup:** reuse structured reference strings (SRS) where applicable.

## 7 Evaluation Methodology

To evaluate an implementation, report:

- prover time per step and end-to-end latency for  $T$  steps;
- peak memory and bandwidth (witness size, transcript size);
- verifier time and proof size for the final proof;
- security parameters (field size, soundness error, transcript model);
- failure modes under adversarial inputs and malformed encodings.

## 8 Related Work

Nova sits within a broader line of work on recursive SNARKs and IVC. Key reference points include:

- recursive composition using pairing-based SNARKs,
- Halo-style recursion without trusted setup,
- SNARK-friendly hash functions and commitment schemes,
- rollups and private computation systems that require efficient recursion.

## 9 Conclusion and Open Problems

Nova-style folding is a compelling foundation for both privacy-preserving protocols and scaling systems. However, production deployments still face challenges in side-channel resistance, robust transcript/version management, curve-cycle availability, and developer ergonomics. Future work includes formal verification of implementations, standardized test vectors, and hardware-aware optimizations.

## References

- [1] Srinath Setty. *Nova: Recursive Zero-Knowledge Arguments from Folding Schemes*. IACR ePrint 2021/370, 2021. <https://eprint.iacr.org/2021/370>
- [2] *CycleFold: Folding on a Cycle of Curves*. IACR ePrint 2023/1192, 2023. <https://eprint.iacr.org/2023/1192>
- [3] Sean Bowe, Jack Grigg, and Daira Hopwood. *Halo: Recursive Proof Composition without a Trusted Setup*. 2019.